

# Topaz

Topaz builds software that learns in real time.

First product: a landing page that optimizes itself. Same substrate, next stop: edge and autonomous systems.

[PENDING-EMMA: confirm one-liner] · [PENDING-A1: A1 page-morphing loop as background]

## THE PROBLEM

# Most websites are too small to A/B test.

- One significant test needs ~400 conversions per variant: about 27,000 visitors at a 3% conversion rate.
- Below that line, the people who make funnels — indie SaaS founders, funnel and landing-page builders — guess, hand-tweak, or freeze whatever an AI page builder gave them.
- Today's "AI optimization" tools route traffic between human-written variants, and gate even that behind \$199-\$249/mo high-traffic tiers. Their statistics need the traffic the customer does not have.

## WHY NOW

# Generate-and-freeze, meet commodity embeddings.

- AI page builders commoditized page creation, and froze page improvement: millions of generated-and-frozen pages have no optimization path.
- The incumbent ceiling is structural (discrete variants), not a tuning gap that will close on its own.
- Frozen embedding spaces became free commodity infrastructure. That is exactly the substrate Sutra compiles onto.

## THE SOLUTION

# The page learns in real time, like an agent.

It improves itself as it runs, not from an outside agent guessing and re-testing it.

- The page IS a program in Sutra, compiled into a continuously adjustable tensor graph.
- Layout, spacing, emphasis, and color are continuous parameters.
- They learn from every scroll, hover, and dwell: dense behavioral signal, not binary conversions.
- So the page improves from the traffic it actually has — and stays readable: "what changed and why" is a real answer, not a heat map.

DEMO

# A1: warmer / colder, the page morphing live.

Demo code complete and containerized (Docker; ~30-min deploy path). [\[PENDING-A1: live URL + recorded video — the two remaining artifacts\]](#)

## EVIDENCE AND THE FOUR RISKS

**We price down risk, then name the next pass/fail milestone.**

**Technical** [retired] Sutra compiler ships (PyPI v0.9.4, five releases since May). NeurIPS submission: bundle decode 100% at k=8, bind/unbind error  $\sim 1.5e-15$ .

**Product** [mostly retired] A working OS prototype in the Topaz stack: terminal, calculator, GUI. The learner beats the best-variant bandit ceiling in simulation (0.824 at 500 visitors vs 0.609).

**Commercial** [open] Customer-discovery outreach wave live (Jul 2026), interviews against pre-written falsifiable hypotheses. Next signal: a structured pilot with named design partners.

**Regulatory** [deferred] None for the web wedge. It enters only with the edge / autonomous-systems markets deliberately sequenced behind it.

**Next milestone (falsifiable):** A2 pre-registered live experiment. Engagement-composite primary metric,  $\sim 250$ -500 post-consent visitors per arm, abort criteria fixed in advance, results published whatever they show. Pass = beats the frozen baseline on real traffic.

## MARKET

# A large, unserved beachhead, with a deep-tech road behind it.

- Beachhead ICP, kept deliberately narrow: the people who MAKE sales funnels — indie SaaS founders and funnel/landing-page builders. Self-serve, card-swipe sales, no procurement.
- Incumbents priced them out: continuous optimization appears only on \$199-\$249/mo high-traffic tiers; sub-500-visitor sites cannot buy it at any sane price. The segment is unserved, not underserved.
- Behind the wedge: edge and autonomous systems, where the whole system being one inspectable network becomes the product. Sequenced, not led with.

**Serviceable beachhead \$57M-\$120M/yr:**  $\geq 165\text{k}-200\text{k}$  paying customers of named funnel tools alone (ClickFunnels 150k disclosed 2023; Unbounce 15k+; Leadpages  $\sim 40\text{k}$ )  $\times$  \$29-\$50/mo. Wider bottom-up TAM  $\sim$  \$7B/yr:  $\sim 24\text{M}$  US small-business sites (SBA 2024)  $\times$   $\sim$  \$300/yr.

Our own arithmetic from public anchors, assumptions stated in the plan; conservative bound uses only company-disclosed paying counts at the \$29 floor.

# The ceiling is structural, not a tuning gap.

- Classical A/B tools need ~27K visitors per significant test. We learn at small-site traffic.
- AI variant routing (Optimizely-class) extracts ~1 bit per visitor and can never beat the best variant a human wrote.
- Topaz optimizes a continuous design space from dense per-visitor signal: a different sample-efficiency regime entirely.



## A language and an OS: the moat and the long game.

- Sutra: a geometrically compiled language; programs become differentiable tensor graphs. On PyPI (sutra-dev v0.9.4), NeurIPS 2026 submission.
- An OS prototype: a GPU-native neuro-symbolic operating system where the whole running system is one inspectable network.
- The page-as-program property, and with it the continuous optimization space, falls out of Sutra's compilation step. A competitor would have to rebuild the substrate to follow.

## TRACTION AND FIRST COMMERCIAL SIGNAL

# What exists, stated plainly.

- Shipped and public: Sutra compiler on PyPI (sutra-dev v0.9.4 — five releases May-Jul 2026), a working OS prototype, NeurIPS 2026 submission.
- In motion: customer-discovery outreach wave sent Jul 2026, interviews against pre-written falsifiable hypotheses; A1 demo code complete and containerized; A2 live experiment pre-registered.
- Pre-product, pre-revenue, no signed pilots yet. The strongest near-term signal we are building toward is a structured pilot with named design partners and written success criteria.

## TEAM

# Solo founder, full-time, operating with AI leverage.

- Emma Leonhart has shipped a working language compiler to PyPI, a working OS prototype (terminal, calculator, click-driven GUI), and a NeurIPS 2026 submission.
- The operating model is AI-leveraged under durable rules, queues, and automated work loops, with the judgment to know when to take the wheel.
- Third-party read, Y Combinator's Paxel analysis: "operating at a strong level... turning messy, multi-repo work into durable rules, releases, queues, cron loops, and truth-preserving engineering constraints."
- The case for backing one founder is this operating record. Hiring multiplies an already-leveraged operator; it does not replace a missing one.

## THE ASK

# Raising to prove the thesis on real traffic.

[PENDING-EMMA: round size + instrument. Tracks: YC + investors, with a cyber.fund-style uncapped SAFE in parallel.]

- Use of funds 1: run the A2 live-traffic experiment at scale (the falsifiable pass/fail milestone above).
- Use of funds 2: land the first structured pilots with named design partners and written success criteria.
- Use of funds 3: relocate operations to San Francisco (already decided).

# **Topaz: software that learns in real time.**

Deep tech with a web wedge that proves the learning loop in public, on a falsifiable schedule — then the same substrate to edge and autonomous systems, behind a language-and-OS moat a competitor would have to rebuild to follow.

Emma Leonhart, Founder · [emma@topazcomputing.com](mailto:emma@topazcomputing.com) · [topazcomputing.com](https://topazcomputing.com)